

Bank Marketing

1. Business Problem

Description

Source: Created by: Paulo Cortez (Univ. Minho) and Sérgio Moro (ISCTE-IUL) @ 2012

<https://archive.ics.uci.edu/ml/datasets/Bank+Marketing>

Dataset: Term - deposit marketing campaign data of a Portuguese banking institution.

Problem Statement: The business problem is a binary classification problem. The classification goal is to predict if the client contacted through the marketing campaign will subscribe a term deposit.

In [1]:

```
# importing requierd libraries
import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
from prettytable import PrettyTable
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import Normalizer
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score
from sklearn.linear_model import SGDClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import roc_curve
from sklearn.metrics import log_loss
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
data = pd.read_csv('bank-full.csv', sep=';')
print('Shape of our data {}'.format(data.shape))
```

Shape of our data (45211, 17)

In [3]:

```
data.head()
```

Out[3]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0

Dataset Description

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be (or not) subscribed.

Attribute/Features Description:

Dataset have 17 attributes including one dependent attribute and there are 45211 instances/datapoints. So we have 16 predictor/independent attributes and 1 dependent attribute.

- **bank client attributes :**
 - age: age of client (numeric)
 - job : type of job (categorical: "admin.", "unknown", "unemployed", "management", "housemaid", "entrepreneur", "student", "blue-collar", "self-employed", "retired", "technician", "services")
 - marital : marital status (categorical: "married", "divorced", "single")
 - education: client highest education (categorical: "unknown", "secondary", "primary", "tertiary")
 - default: has credit in default? (binary/2-categories: "yes", "no")
 - balance: average yearly balance, in euros (numeric)
 - housing: has housing loan? (binary/2-categories: "yes", "no")
 - loan: has personal loan? (binary/2-categories: "yes", "no")
- **related with the last contact of the current campaign:**
 - contact: contact communication type (categorical: "unknown", "telephone", "cellular")
 - day: last contact day of the month (numeric)
 - month: last contact month of year (categorical: "jan", "feb", "mar", ..., "nov", "dec")
 - duration: last contact duration, in seconds (numeric)
- **other attributes:**
 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)
 - previous: number of contacts performed before this campaign and for this client (numeric)
 - poutcome: outcome of the previous marketing campaign (categorical: 'unknown',"other", "failure", "success")
- **Output variable (desired target):**
 - y: has the client subscribed a term deposit? (binary: "yes", "no")

In [4]:

```
data.describe(include='all')
```

Out[4]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration
count	45211.000000	45211	45211	45211	45211	45211.000000	45211	45211	45211	45211.000000	45211	45211.000000
unique	NaN	12	3	4	2	NaN	2	2	3	NaN	12	NaN
top	NaN	blue-collar	married	secondary	no	NaN	yes	no	cellular	NaN	may	NaN
freq	NaN	9732	27214	23202	44396	NaN	25130	37967	29285	NaN	13766	NaN
mean	40.936210	NaN	NaN	NaN	NaN	1362.272058	NaN	NaN	NaN	15.806419	NaN	258.163080
std	10.618762	NaN	NaN	NaN	NaN	3044.765829	NaN	NaN	NaN	8.322476	NaN	257.527812
min	18.000000	NaN	NaN	NaN	NaN	-8019.000000	NaN	NaN	NaN	1.000000	NaN	0.000000
25%	33.000000	NaN	NaN	NaN	NaN	72.000000	NaN	NaN	NaN	8.000000	NaN	103.000000
50%	39.000000	NaN	NaN	NaN	NaN	448.000000	NaN	NaN	NaN	16.000000	NaN	180.000000
75%	48.000000	NaN	NaN	NaN	NaN	1428.000000	NaN	NaN	NaN	21.000000	NaN	319.000000
max	95.000000	NaN	NaN	NaN	NaN	102127.000000	NaN	NaN	NaN	31.000000	NaN	4918.000000

In [5]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
age                45211 non-null int64
job                45211 non-null object
```

```

job            45211 non-null object
marital        45211 non-null object
education      45211 non-null object
default        45211 non-null object
balance        45211 non-null int64
housing        45211 non-null object
loan           45211 non-null object
contact        45211 non-null object
day            45211 non-null int64
month          45211 non-null object
duration       45211 non-null int64
campaign       45211 non-null int64
pdays        45211 non-null int64
previous       45211 non-null int64
poutcome      45211 non-null object
y             45211 non-null object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB

```

Observation:

Our dataset do not have any null/nan/missing values.

In [6]:

```

categorical = ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome']
numerical = [x for x in data.columns.to_list() if x not in categorical]
numerical.remove('y')

```

In [7]:

```

print('Categorical features:', categorical)
print('Numerical features:', numerical)

```

```

Categorical features: ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome']
Numerical features: ['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']

```

In [8]:


```

from matplotlib import pyplot as plt
sns.countplot(x=data['y'])
plt.title('Distribution of classes')
plt.xlabel('Target class')

```

Out[8]:

Text(0.5, 0, 'Target class')



In [9]:

```
data.y.value_counts()
```

Out[9]:

```
no      39922
yes      5289
Name: y, dtype: int64
```

Observation:

Our dataset is highly imbalanced.

Data Analysis


pdays

In [10]:

```
sns.boxplot(y=data['pdays'], x=data['y'])
plt.title('Box plot of pdays vs y (target variable)')
plt.xlabel('y: target variable')
```

Out[10]:

Text(0.5, 0, 'y: target variable')





In [11]:

```
sns.FacetGrid(data, hue='y', size=5) \
.map(sns.distplot, 'pdays') \
.add_legend()
plt.title('PDF of pdays for target variable y')
```

Out[11]:

Text(0.5, 1, 'PDF of pdays for target variable y')





In [12]:

```
data.pdays.describe()
```

Out[12]:

```
count    45211.000000
mean      40.197828
std       100.128746
min       -1.000000
25%       -1.000000
50%       -1.000000
75%       -1.000000
max       871.000000
Name: pdays, dtype: float64
```

In [13]:

```
for x in range(95, 101, 1):
    print("{}% of pdays are less than equal to {}".format(x, data.pdays.quantile(x/100)))
iqr = data.pdays.quantile(0.75) - data.pdays.quantile(0.25)
print('IQR {}'.format(iqr))
```

```
95% of pdays are less than equal to 317.0
96% of pdays are less than equal to 337.0
97% of pdays are less than equal to 349.0
98% of pdays are less than equal to 360.0
99% of pdays are less than equal to 370.0
100% of pdays are less than equal to 871.0
IQR 0.0
```

Observation:

- The attribute pdays seems to be important feature as there is a clear distinction in quartile ranges of pdays for target variable yes and no.
- 75% clients contacted through campaign are not previously contacted.
- Mean of pdays is 40.20
- There are outliers as we can see from boxplot.

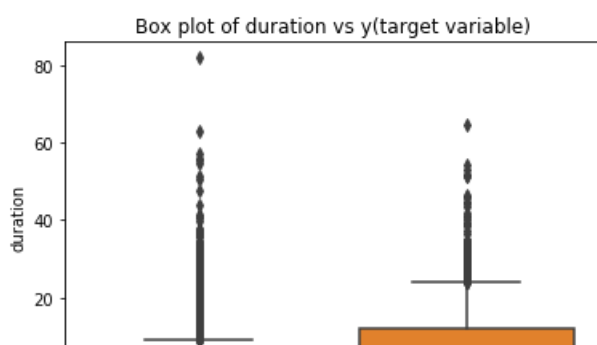
duration

In [14]:

```
# converting call duration from seconds to minute
data['duration'] = data['duration']/60
sns.boxplot(y=data['duration'], x=data['y'])
plt.title('Box plot of duration vs y(target variable)')
plt.xlabel('y:target variable')
```

Out[14]:

```
Text(0.5, 0, 'y:target variable')
```






In [17]:

```
sns.FacetGrid(data, hue='y', size=5) \
.map(sns.distplot, 'duration') \
.add_legend()
plt.title('PDF of duration for target variable y')
```

Out[17]:

Text(0.5, 1, 'PDF of duration for target variable y')



In [18]:

```
data.duration.describe()
```

Out[18]:

```
count    45211.000000
mean       4.302718
std        4.292130
min         0.000000
25%        1.716667
50%        3.000000
75%        5.316667
max       81.966667
Name: duration, dtype: float64
```

In [19]:

```
for x in range(95, 101, 1):
    print("{}% of calls have duration less than equal to {}".format(x, data.duration.quantile(x/100)))
iqr = data.duration.quantile(0.75) - data.duration.quantile(0.25)
print('IQR {}'.format(iqr))
```

```
95% of calls have duration less than equal to 12.516666666666667
96% of calls have duration less than equal to 13.716666666666667
97% of calls have duration less than equal to 15.244999999999951
98% of calls have duration less than equal to 17.516666666666666
99% of calls have duration less than equal to 21.15
100% of calls have duration less than equal to 81.96666666666667
IQR 3.5999999999999996
```

Observation:

Observation:

- The attribute duration seems to be important feature as there is a clear distinction in quartile ranges of duration for target variable yes and no.
- 75% call duration are less than or equal to 5.32
- duration have a mean of 4.30 and standard-deviation 4.29
- There are outliers points in duration.

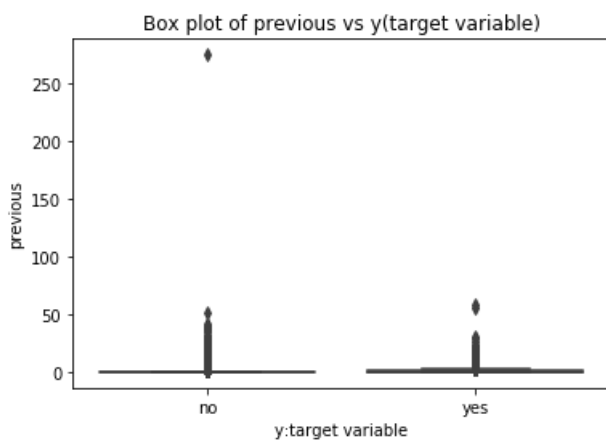
previous

In [20]:

```
sns.boxplot(y=data['previous'], x=data['y'])  
plt.title('Box plot of previous vs y(target variable)')  
plt.xlabel('y:target variable')
```

Out[20]:

Text(0.5, 0, 'y:target variable')




In [21]:

```
sns.FacetGrid(data, hue='y', size=5) \  
.map(sns.distplot, 'previous') \  
.add_legend()  
plt.title('PDF of previous values for target variable y')
```

Out[21]:

Text(0.5, 1, 'PDF of previous values for target variable y')



In [22]:

```
data.previous.describe()
```

Out[22]:

```
count      45211.000000
mean         0.580323
std          2.303441
min           0.000000
25%           0.000000
50%           0.000000
75%           0.000000
max          275.000000
Name: previous, dtype: float64
```

In [23]:

```
for x in range(95, 101, 1):
    print("{}% of previous values less than equal to {}".format(x, data.previous.quantile(x/100)))
iqr = data.previous.quantile(0.75) - data.previous.quantile(0.25)
print('IQR {}'.format(iqr))
```

```
95% of previous values less than equal to 3.0
96% of previous values less than equal to 4.0
97% of previous values less than equal to 5.0
98% of previous values less than equal to 6.0
99% of previous values less than equal to 8.9000000000001455
100% of previous values less than equal to 275.0
IQR 0.0
```

Observation:

- 75% of previous values equal 0 and 99% values ≤ 8.90
- duration have a mean of 0.58 and standard-deviation 2.30
- There are outliers points in duration.


age

In [24]:

```
sns.boxplot(y=data['age'], x=data['y'])
plt.title('Box plot of age vs y(target variable)')
plt.xlabel('y:target variable')
```

Out[24]:

Text(0.5, 0, 'y:target variable')



In [25]:


```
sns.FacetGrid(data, hue='y', size=5) \
    map(sns.distplot, 'age') \
```



```
map sns.distplot, age, \
.add_legend()
plt.title('PDF of age for target variable y')
```

Out[25]:

Text(0.5, 1, 'PDF of age for target variable y')



In [26]:

```
data.age.describe()
```

Out[26]:

```
count    45211.000000
mean      40.936210
std       10.618762
min       18.000000
25%       33.000000
50%       39.000000
75%       48.000000
max       95.000000
Name: age, dtype: float64
```

In [27]:

```
for x in range(95, 101, 1):
    print("{}% of people having age are less than equal to {}".format(x, data.age.quantile(x/100)))
iqr = data.age.quantile(0.75) - data.age.quantile(0.25)
print('IQR {}'.format(iqr))
```

```
95% of people having age are less than equal to 59.0
96% of people having age are less than equal to 59.0
97% of people having age are less than equal to 60.0
98% of people having age are less than equal to 63.0
99% of people having age are less than equal to 71.0
100% of people having age are less than equal to 95.0
IQR 15.0
```

In [28]:

```
lst = [data]
for column in lst:
    column.loc[column["age"] < 30, 'age_group'] = 30
    column.loc[(column["age"] >= 30) & (column["age"] <= 44), 'age_group'] = 40
    column.loc[(column["age"] >= 45) & (column["age"] <= 59), 'age_group'] = 50
    column.loc[column["age"] >= 60, 'age_group'] = 60
```

In [29]:


```
count_age_response_pct = pd.crosstab(data['y'], data['age_group']).apply(lambda x: x/x.sum() * 100)
count_age_response_pct = count_age_response_pct.transpose()
```

In [30]:

```
sns.countplot(x='age_group', data=data, hue='y')
```

Out[30]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f2574681d68>



In [31]:

```
print('Success rate and total clients contacted for different age_groups:')
print('Clients age < 30 contacted: {}, Success rate: {}'.format(len(data[data['age_group'] == 30]),
data[data['age_group'] == 30].y.value_counts()[1]/len(data[data['age_group'] == 30])))
print('Clients of age 30-45 contacted: {}, Success rate: {}'.format(len(data[data['age_group'] == 40]),
data[data['age_group'] == 40].y.value_counts()[1]/len(data[data['age_group'] == 40])))
print('Clients of age 40-60 contacted: {}, Success rate: {}'.format(len(data[data['age_group'] == 50]),
data[data['age_group'] == 50].y.value_counts()[1]/len(data[data['age_group'] == 50])))
print('Clients of 60+ age contacted: {}, Success rate: {}'.format(len(data[data['age_group'] == 60]),
data[data['age_group'] == 60].y.value_counts()[1]/len(data[data['age_group'] == 60])))
```

```
Success rate and total clients contacted for different age_groups:
Clients age < 30 contacted: 5273, Success rate: 0.1759908970225678
Clients of age 30-45 contacted: 24274, Success rate: 0.10117821537447474
Clients of age 40-60 contacted: 13880, Success rate: 0.09402017291066282
Clients of 60+ age contacted: 1784, Success rate: 0.336322869955157
```

Observation:

- People with age < 30 or 60+ have higher success rate.
- Only 3% of clients have age of 60+

jobs

In [32]:

```
data.job.value_counts()
```

Out[32]:

```
blue-collar    9732
management    9458
technician     7597
admin.         5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student        938
```


```
student          938
unknown          288
Name: job, dtype: int64
```

In [33]:

```
sns.set(rc={'figure.figsize':(20,6)})
sns.countplot(x=data['job'], data=data, hue=data['y'])
plt.title('Count Plot of job for target variable')
```

Out[33]:

Text(0.5, 1.0, 'Count Plot of job for target variable')



In [34]:

```
table = PrettyTable(['Job', 'Total Clients', 'Success rate'])
table.add_row(['Blue-collar', len(data[data['job'] == 'blue-collar']), data[data['job'] == 'blue-collar'].y.value_counts()[1]/len(data[data['job'] == 'blue-collar'])])
table.add_row(['Management', len(data[data['job'] == 'management']), data[data['job'] == 'management'].y.value_counts()[1]/len(data[data['job'] == 'management'])])
table.add_row(['Technician', len(data[data['job'] == 'technician']), data[data['job'] == 'technician'].y.value_counts()[1]/len(data[data['job'] == 'technician'])])
table.add_row(['Admin', len(data[data['job'] == 'admin.']), data[data['job'] == 'admin.'].y.value_counts()[1]/len(data[data['job'] == 'admin.'])])
table.add_row(['Services', len(data[data['job'] == 'services']), data[data['job'] == 'services'].y.value_counts()[1]/len(data[data['job'] == 'services'])])
table.add_row(['Retired', len(data[data['job'] == 'retired']), data[data['job'] == 'retired'].y.value_counts()[1]/len(data[data['job'] == 'retired'])])
table.add_row(['Self-employed', len(data[data['job'] == 'self-employed']), data[data['job'] == 'self-employed'].y.value_counts()[1]/len(data[data['job'] == 'self-employed'])])
table.add_row(['Entrepreneur', len(data[data['job'] == 'entrepreneur']), data[data['job'] == 'entrepreneur'].y.value_counts()[1]/len(data[data['job'] == 'entrepreneur'])])
table.add_row(['Unemployed', len(data[data['job'] == 'unemployed']), data[data['job'] == 'unemployed'].y.value_counts()[1]/len(data[data['job'] == 'unemployed'])])
table.add_row(['Housemaid', len(data[data['job'] == 'housemaid']), data[data['job'] == 'housemaid'].y.value_counts()[1]/len(data[data['job'] == 'housemaid'])])
table.add_row(['Student', len(data[data['job'] == 'student']), data[data['job'] == 'student'].y.value_counts()[1]/len(data[data['job'] == 'student'])])
table.add_row(['Unknown', len(data[data['job'] == 'unknown']), data[data['job'] == 'unknown'].y.value_counts()[1]/len(data[data['job'] == 'unknown'])])
print(table)
```

Job	Total Clients	Success rate
Blue-collar	9732	0.07274969173859433
Management	9458	0.13755550856417847
Technician	7597	0.11056996182703699
Admin	5171	0.12202668729452718
Services	4154	0.08883004333172845
Retired	2264	0.22791519434628976
Self-employed	1579	0.11842938568714376
Entrepreneur	1487	0.08271687962340282
Unemployed	1303	0.15502686108979277
Housemaid	1240	0.08790322580645162
Student	938	0.2867803837953092

```
|      Unknown      |      288      | 0.11805555555555555 |
+-----+-----+-----+
```

Observation:

- Top contacted clients are from job type: 'blue-collar', 'management' & 'technician'
- Success rate is highest for student

poutcome

In [35]:

```
data.poutcome.value_counts()
```

Out[35]:


```
unknown    36959
failure     4901
other       1840
success     1511
Name: poutcome, dtype: int64
```

In [36]:

```
sns.countplot(x=data['poutcome'], data=data, hue=data['y'])
plt.title('Count Plot of poutcome for target variable')
```

Out[36]:

Text(0.5, 1.0, 'Count Plot of poutcome for target variable')



Observation:

- Most of the clients contacted have previous outcome as 'unknown'.

education

In [37]:

```
data.education.value_counts()
```

Out[37]:


```
secondary    23202
tertiary     13301
primary       6851
unknown       1857
Name: education, dtype: int64
```

In [38]:

```
sns.countplot(x=data['education'], data=data, hue=data['y'])
plt.title('Count plot of education for target variable y')
```

Out[38]:

Text(0.5, 1.0, 'Count plot of education for target variable y')



Observation:

- Most of the people who are contacted have tertiary or secondary education.

default

In [39]:

```
data.default.value_counts()
```

Out[39]:

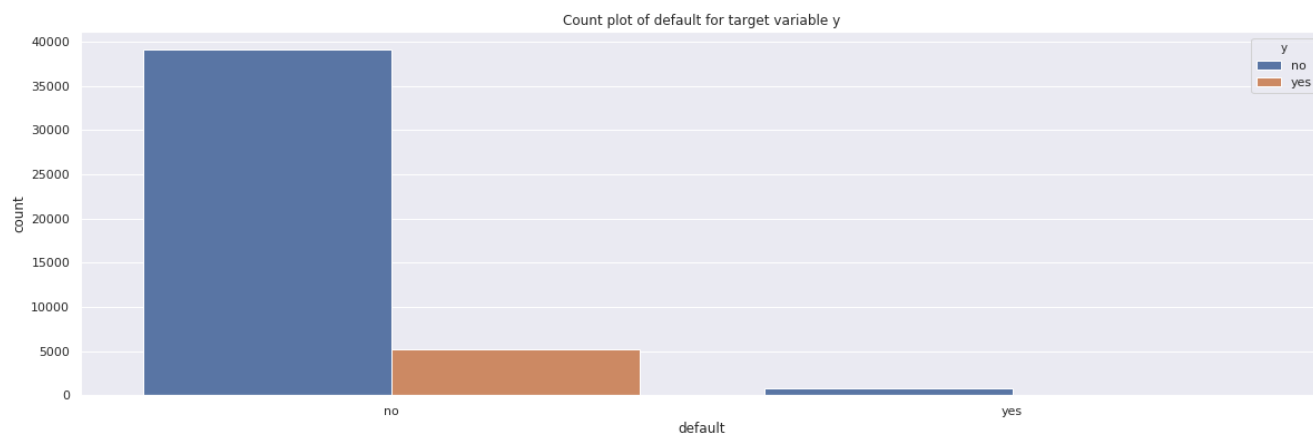
```
no      44396
yes       815
Name: default, dtype: int64
```

In [40]:

```
sns.countplot(x=data['default'], data=data, hue=data['y'])
plt.title('Count plot of default for target variable y')
```

Out[40]:

Text(0.5, 1.0, 'Count plot of default for target variable y')



In [41]:

```
data[data['default'] == 'yes'].y.count()
```

Out[41]:

815

Observation:

Very few clients are contacted who are defaulter,

loan

In [42]:

```
data.loan.value_counts()
```

Out[42]:


```
no      37967
yes      7244
Name: loan, dtype: int64
```

In [43]:

```
sns.countplot(x=data['loan'], data=data, hue=data['y'])
plt.title('Count plot of loan for target variable y')
```

Out[43]:

Text(0.5, 1.0, 'Count plot of loan for target variable y')



Observation:

- As seen for default variable, less client are contacted who have loan.

contact

In [44]:

```
data.contact.value_counts()
```

Out[44]:

```
cellular      29285
unknown       13020
telephone      2906
Name: contact, dtype: int64
```


In [45]:

```
sns.countplot(x=data['contact'], data=data, hue=data['y'])
```

```
plt.title('Count plot of contact for target variable y')
```

Out[45]:

Text(0.5, 1.0, 'Count plot of contact for target variable y')



Observation:

Most of the people are contacted through cellular

month

In [46]:

```
data.month.value_counts()
```

Out[46]:


```
may    13766
jul     6895
aug     6247
jun     5341
nov     3970
apr     2932
feb     2649
jan     1403
oct       738
sep       579
mar       477
dec       214
Name: month, dtype: int64
```

In [47]:

```
sns.countplot(x=data['month'], data=data, hue=data['y'])
plt.title('Count plot of month for target variable y')
```

Out[47]:

Text(0.5, 1.0, 'Count plot of month for target variable y')





In [48]:

```
data[data['month'] == 'jan'].y.value_counts()
```

Out[48]:

```
no      1261
yes      142
Name: y, dtype: int64
```

In [49]:

```
print('Success rate and total clients contacted for different months:')
print('Clients contacted in January: {}, Success rate: {}'.format(len(data[data['month'] ==
'jan']), data[data['month'] == 'jan'].y.value_counts()[1]/len(data[data['month'] == 'jan'])))
print('Clients contacted in February: {}, Success rate: {}'.format(len(data[data['month'] == 'feb'
]), data[data['month'] == 'feb'].y.value_counts()[1]/len(data[data['month'] == 'feb'])))
print('Clients contacted in March: {}, Success rate: {}'.format(len(data[data['month'] == 'mar'
]), data[data['month'] == 'mar'].y.value_counts()[1]/len(data[data['month'] == 'mar'])))
print('Clients contacted in April: {}, Success rate: {}'.format(len(data[data['month'] == 'apr'
]), data[data['month'] == 'apr'].y.value_counts()[1]/len(data[data['month'] == 'apr'])))
print('Clients contacted in May: {}, Success rate: {}'.format(len(data[data['month'] == 'may'
]), data[data['month'] == 'may'].y.value_counts()[1]/len(data[data['month'] == 'may'])))
print('Clients contacted in June: {}, Success rate: {}'.format(len(data[data['month'] == 'jun'
]), data[data['month'] == 'jun'].y.value_counts()[1]/len(data[data['month'] == 'jun'])))
print('Clients contacted in July: {}, Success rate: {}'.format(len(data[data['month'] == 'jul'
]), data[data['month'] == 'jul'].y.value_counts()[1]/len(data[data['month'] == 'jul'])))
print('Clients contacted in August: {}, Success rate: {}'.format(len(data[data['month'] == 'aug'
]), data[data['month'] == 'aug'].y.value_counts()[1]/len(data[data['month'] == 'aug'])))
print('Clients contacted in September: {}, Success rate: {}'.format(len(data[data['month'] == 'sep'
]), data[data['month'] == 'sep'].y.value_counts()[1]/len(data[data['month'] == 'sep'])))
print('Clients contacted in October: {}, Success rate: {}'.format(len(data[data['month'] ==
'oct']), data[data['month'] == 'oct'].y.value_counts()[1]/len(data[data['month'] == 'oct'])))
print('Clients contacted in November: {}, Success rate: {}'.format(len(data[data['month'] == 'nov'
]), data[data['month'] == 'nov'].y.value_counts()[1]/len(data[data['month'] == 'nov'])))
print('Clients contacted in December: {}, Success rate: {}'.format(len(data[data['month'] == 'dec'
]), data[data['month'] == 'dec'].y.value_counts()[1]/len(data[data['month'] == 'dec'])))
```

```
Success rate and total clients contacted for different months:
Clients contacted in January: 1403, Success rate: 0.10121168923734854
Clients contacted in February: 2649, Success rate: 0.1664779161947905
Clients contacted in March: 477, Success rate: 0.480083857442348
Clients contacted in April: 2932, Success rate: 0.19679399727148705
Clients contacted in May: 13766, Success rate: 0.06719453726572715
Clients contacted in June: 5341, Success rate: 0.10222804718217562
Clients contacted in July: 6895, Success rate: 0.09093546047860769
Clients contacted in August: 6247, Success rate: 0.11013286377461182
Clients contacted in September: 579, Success rate: 0.46459412780656306
Clients contacted in October: 738, Success rate: 0.43766937669376693
Clients contacted in November: 3970, Success rate: 0.10151133501259446
Clients contacted in December: 214, Success rate: 0.4672897196261682
```

Observation:

- Most of the clients (approx 1/3 of total) are contacted in the month of May but the success rate is only 6.7%.
- March have highest success rate.

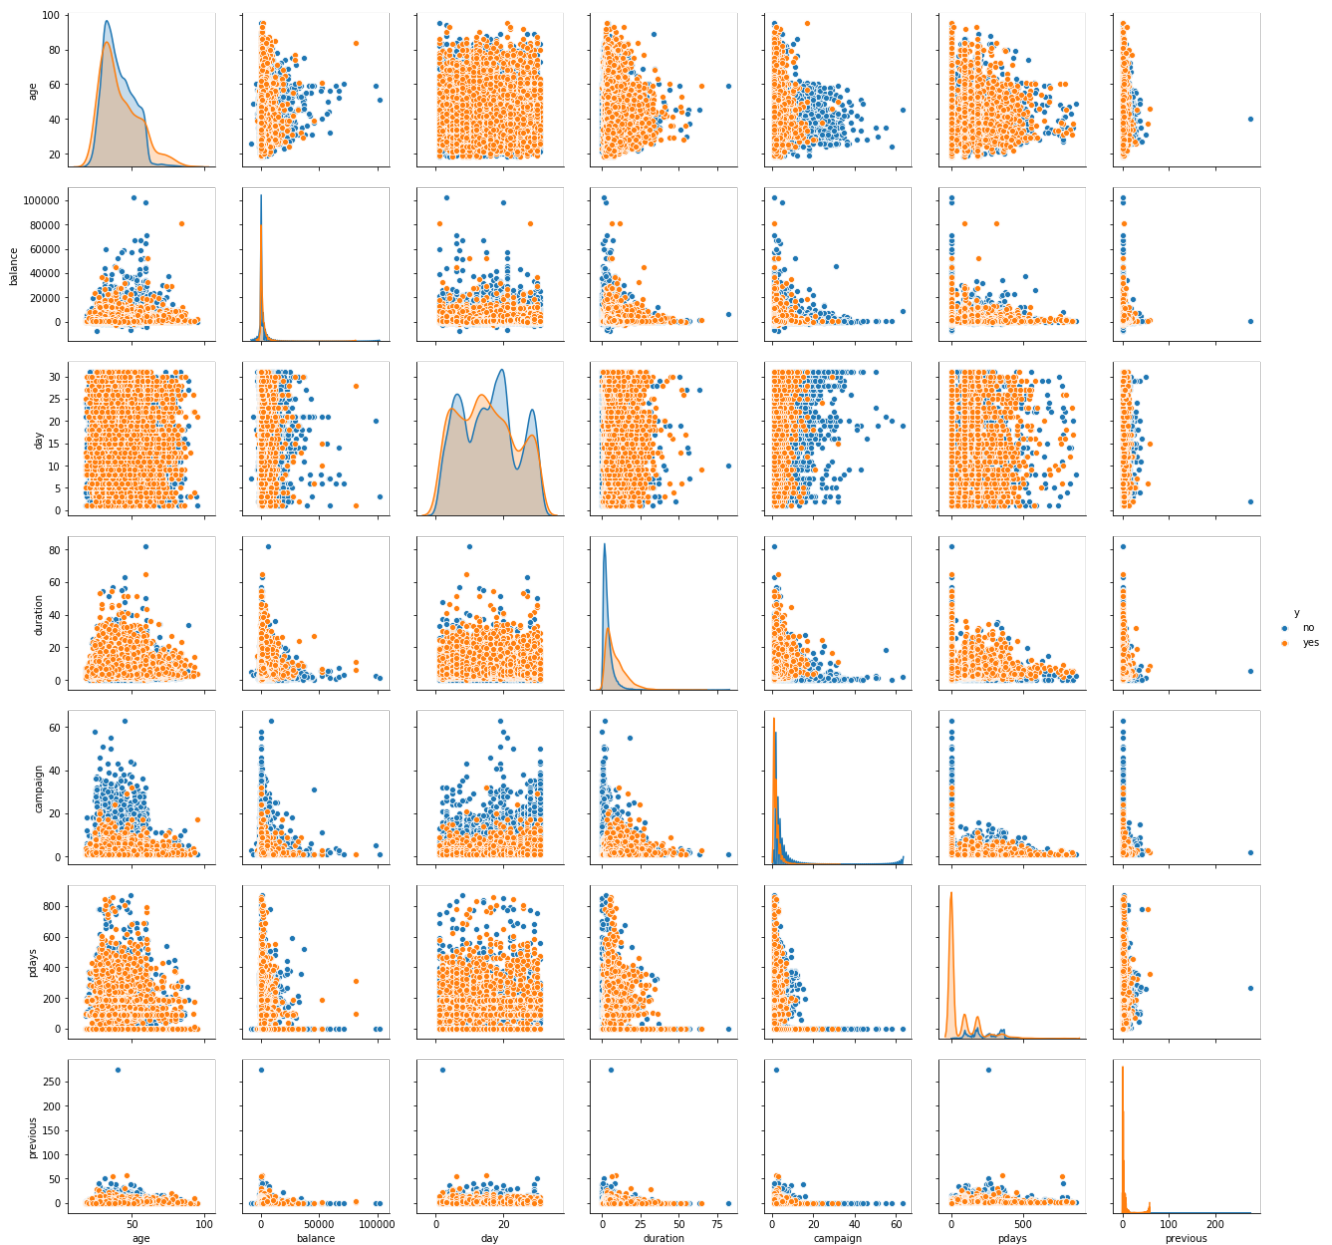
Pairplot

In [15]:

```
#data.drop('age_group', axis=1, inplace=True)
sns.pairplot(data, hue='y')
```

Out[15]:

```
<seaborn.axisgrid.PairGrid at 0x7facfc4866a0>
```

Observation:

- For most of the variables our pair plot is overlapping a lot.
- Pair plots of age-campaign and day-campaign are much efficient in distinguishing between different classes with very few overlaps.

Correlation matrix of numerical features

In [17]:

```
corr_data = data[numerical + ['y']]
corr = corr_data.corr()
plt.close()
cor_plot = sns.heatmap(corr, annot=True, cmap='RdYlGn', linewidths=0.2, annot_kws={'size':10})
fig=plt.gcf()
fig.set_size_inches(12,10)
plt.xticks(fontsize=10,rotation=-30)
plt.yticks(fontsize=10)
plt.title('Correlation Matrix')
plt.show()
```





Observation:

- Over numerical features have very less correlation between them.
- pdays and previous have higher correlation
- duration have a higher correlation with our target variable

Outlier detection for numerical attributes using IQR

In [18]:

```
# creating new data frame of numerical columns
data_numerical = data[numerical]
print('Shape of numerical dataframe {}'.format(data_numerical.shape))
data_numerical.head()
```

Shape of numerical dataframe (45211, 7)

Out[18]:

	age	balance	day	duration	campaign	pdays	previous
0	58	2143	5	4.350000	1	-1	0
1	44	29	5	2.516667	1	-1	0
2	33	2	5	1.266667	1	-1	0
3	47	1506	5	1.533333	1	-1	0
4	33	1	5	3.300000	1	-1	0

In [19]:

```
q3 = data_numerical.quantile(0.75)
q1 = data_numerical.quantile(0.25)
iqr = q3 - q1
print('IQR for numerical attributes')
print(iqr)
```

```
IQR for numerical attributes
age          15.0
balance      1356.0
day          13.0
duration     3.6
campaign     2.0
pdays       0.0
previous     0.0
dtype: float64
```

In [20]:

```
data_out = data[~((data_numerical < (q1 - 1.5 * iqr)) | (data_numerical > (q3 + 1.5 *
iqr))).any(axis=1)]
print('{} points are outliers based on IQR'.format(data.shape[0] - data_out.shape[0]))
```

17029 points are outliers based on IQR

In [21]:

```
data.shape
```

Out[21]:

(45211, 17)

Preprocessing

Train Test Split

In [3]:

```
data.replace(to_replace={'y':'yes'}, value=1, inplace=True)
data.replace(to_replace={'y':'no'}, value=0, inplace=True)
```

In [22]:

```
# Convert the columns into categorical variables
data1 = data.copy()
data1['job'] = data1['job'].astype('category').cat.codes
data1['marital'] = data1['marital'].astype('category').cat.codes
data1['education'] = data1['education'].astype('category').cat.codes
data1['contact'] = data1['contact'].astype('category').cat.codes
data1['poutcome'] = data1['poutcome'].astype('category').cat.codes
data1['month'] = data1['month'].astype('category').cat.codes
data1['default'] = data1['default'].astype('category').cat.codes
data1['loan'] = data1['loan'].astype('category').cat.codes
data1['housing'] = data1['housing'].astype('category').cat.codes
```

In [4]:

```
y = data['y']
x_train, x_test, y_train, y_test = train_test_split(data.drop(['y'], axis=1), y, test_size=0.20, random_state=42)
```

In [24]:

```
print('Train data shape {} {}'.format(x_train.shape, y_train.shape))
print('Test data shape {} {}'.format(x_test.shape, y_test.shape))
```

Train data shape (36168, 16) (36168,)
Test data shape (9043, 16) (9043,)

Feature Importance

In [25]:


```
plt.style.use('seaborn-white')

clf = DecisionTreeClassifier(class_weight='balanced', min_weight_fraction_leaf = 0.01)

clf.fit(x_train, y_train)
importances = clf.feature_importances_
feature_names = data.drop('y', axis=1).columns
indices = np.argsort(importances)

def feature_importance_graph(indices, importances, feature_names):
    plt.figure(figsize=(12,6))
    plt.title("Feature Importance", fontsize=10)
    plt.barh(range(len(indices)), importances[indices], color='g', align="center")
    plt.yticks(range(len(indices)), feature_names[indices], rotation='horizontal', fontsize=14)
    plt.ylim([-1, len(indices)])

feature_importance_graph(indices, importances, feature_names)
plt.show()
```



Important features we are going to consider for machine learning models:

- duration
- contact
- poutcome
- month
- housing
- pdays
- age
- balance

Encoding data

Encoding categories

In [5]:

```
vectorizer = CountVectorizer(vocabulary=x_train.poutcome.unique())
x_train_poutcome = vectorizer.fit_transform(x_train.poutcome)
x_test_poutcome = vectorizer.transform(x_test.poutcome)
```

In [6]:

```
vectorizer = CountVectorizer(vocabulary=x_train.contact.unique())
```

```
vectorizer = CountVectorizer(vocabulary=x_train.contact.unique())
x_train_contact = vectorizer.fit_transform(x_train.contact)
x_test_contact = vectorizer.transform(x_test.contact)
```

In [7]:

```
vectorizer = CountVectorizer(vocabulary=x_train.month.unique())
x_train_month = vectorizer.fit_transform(x_train.month)
x_test_month = vectorizer.transform(x_test.month)
```

In [8]:

```
vectorizer = CountVectorizer(vocabulary=x_train.housing.unique())
x_train_housing = vectorizer.fit_transform(x_train.housing)
x_test_housing = vectorizer.transform(x_test.housing)
```

Encoding Numerical data using Normalizer()

In [9]:

```
vectorizer = Normalizer()
x_train_duration = vectorizer.fit_transform(x_train.duration.values.reshape(1,-1)).transpose()
x_test_duration = vectorizer.transform(x_test.duration.values.reshape(1, -1)).transpose()
```

In [10]:

```
vectorizer = Normalizer()
x_train_pdays = vectorizer.fit_transform(x_train.pdays.values.reshape(1,-1)).transpose()
x_test_pdays = vectorizer.transform(x_test.pdays.values.reshape(1, -1)).transpose()
```

In [11]:

```
vectorizer = Normalizer()
x_train_age = vectorizer.fit_transform(x_train.age.values.reshape(1,-1)).transpose()
x_test_age = vectorizer.transform(x_test.age.values.reshape(1, -1)).transpose()
```

In [12]:

```
vectorizer = Normalizer()
x_train_balance = vectorizer.fit_transform(x_train.balance.values.reshape(1,-1)).transpose()
x_test_balance = vectorizer.transform(x_test.balance.values.reshape(1, -1)).transpose()
```

In [13]:

```
from scipy.sparse import hstack

train = hstack((x_train_contact, x_train.poutcome, x_train_month, x_train_housing, x_train_duration,
, x_train_pdays, x_train_age, x_train_balance)).tocsr()

test = hstack((x_test_contact, x_test.poutcome, x_test_month, x_test_housing, x_test_duration, x_test_pdays,
x_test_age, x_test_balance)).tocsr()
```

Machine Learning Models

In [15]:

```
# dictionary to store accuracy and roc score for each model
score = {}
```

Logistic Regression

Hyperparameter tuning Logistic Regression

In [37]:

```
parameters = {'C': [(10**i)*x for i in range(-4, 1) for x in [1,3,5]]}

model = LogisticRegression(class_weight='balanced')
clf = RandomizedSearchCV(model, parameters, cv=5, scoring='roc_auc', return_train_score=True, n_jobs=-1)
clf.fit(train, y_train)
print('Best parameters: {}'.format(clf.best_params_))
print('Best score: {}'.format(clf.best_score_))
```

Best parameters: {'C': 3}
Best score: 0.8496734277981354

Training Logistic Regression with best hyperparameters

In [16]:

```
from sklearn.metrics import log_loss

model = LogisticRegression(C=3, class_weight='balanced', n_jobs=-1)
model.fit(train, y_train)
y_probs_train = model.predict_proba(train)
y_probs_test = model.predict_proba(test)
y_predicted_train = model.predict(train)
y_predicted_test = model.predict(test)


# keep probabilities for the positive outcome only
y_probs_train = y_probs_train[:, 1]
y_probs_test = y_probs_test[:, 1]


# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['Logistic Regression'] = [test_auc, test_acc]

# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()
```

```
*****
Train AUC: 0.859
Test AUC: 0.887
*****
Train Accuracy: 0.833
Test Accuracy: 0.746
```





Train Confusion Matrix

In [19]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[27108  4862]
 [ 1172  3026]]
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3acc186438>



Test Confusion Matrix

In [20]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[5795 2157]
 [ 136  955]]
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3acc4a3ba8>



Random Forest

Hyperparameter tuning Random Forest

In [40]:

```
params = {'n_estimators':[75, 100, 250, 500], 'max_depth':[3, 5, 10, 15, 25]}
model = RandomForestClassifier(class_weight='balanced', n_jobs=-1)
clf = RandomizedSearchCV(model, param_distributions=params, cv=5, scoring='roc_auc', random_state=42, n_jobs=-1, return_train_score=True)
clf.fit(train, y_train)
print('Best parameters: {}'.format(clf.best_params_))
print('Best score: {}'.format(clf.best_score_))
```

Best parameters: {'n_estimators': 250, 'max_depth': 25}
Best score: 0.918424480262767

Training random forest with best hyperparameters

In [21]:

```
model = RandomForestClassifier(n_estimators=250, max_depth=25, class_weight='balanced', n_jobs=-1)
model.fit(train, y_train)
y_probs_train = model.predict_proba(train)
y_probs_test = model.predict_proba(test)
y_predicted_train = model.predict(train)
y_predicted_test = model.predict(test)

# keep probabilities for the positive outcome only
y_probs_train = y_probs_train[:, 1]
y_probs_test = y_probs_test[:, 1]

# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
```



```

print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['Random Forest'] = [test_auc, test_acc]


# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()

```

```

*****
Train AUC: 1.000
Test AUC: 0.799
*****
Train Accuracy: 0.999
Test Accuracy: 0.812

```



Train Confusion Matrix

In [22]:

```

from sklearn.metrics import confusion_matrix

cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')

```

Confusion matrix:


```


[[31944  26]
 [ 1  4197]]

```

Out[22]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3acc1a0ef0>





Test Confusion Matrix

In [24]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[6815 1137]
 [ 561  530]]
```

Out[24]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3abc9e4cf8>



SVM

Hyperparameter tuning SVM

In [43]:

```
params = {'alpha': [10**i for i in range(-4, 5)]}

model = SGDClassifier(class_weight='balanced', n_jobs=-1)
```

```

clf = RandomizedSearchCV(model, param_distributions=params, cv=5, scoring='roc_auc', random_state=42, n_jobs=-1, return_train_score=True)
clf.fit(train, y_train)
print('Best parameters: {}'.format(clf.best_params_))
print('Best score: {}'.format(clf.best_score_))

```

Best parameters: {'alpha': 0.0001}
Best score: 0.7767774627832619

Training SVM with best hyperparameters

In [25]:

```

model = SGDClassifier(alpha=0.0001, class_weight='balanced', n_jobs=-1)
model.fit(train, y_train)
y_probs_train = model.decision_function(train)
y_probs_test = model.decision_function(test)
y_predicted_train = model.predict(train)
y_predicted_test = model.predict(test)

# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['SVM'] = [test_auc, test_acc]


# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()

```

```

*****
Train AUC: 0.776
Test AUC: 0.808
*****
Train Accuracy: 0.809
Test Accuracy: 0.805

```



Train Confusion Matrix

In [26]:


```
from sklearn.metrics import confusion_matrix

cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:
[[27132 4838]
 [2056 2142]]

Out[26]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3abc909a90>



Test Confusion Matrix

In [27]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:
[[6698 1254]
 [506 585]]

Out[27]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3abc83ac18>





XGBoost

Hyperparameter tuning XGBClassifier

In [45]:

```
from xgboost import XGBClassifier

params = {'max_depth': [5, 10, 15], 'n_estimators': [10, 100, 500]}

model = XGBClassifier(class_weight='balanced', n_jobs=-1)
clf = RandomizedSearchCV(model, param_distributions=params, cv=5, scoring='roc_auc', random_state=42, n_jobs=-1, return_train_score=True)
clf.fit(train, y_train)
print('Best parameters: {}'.format(clf.best_params_))
print('Best score: {}'.format(clf.best_score_))
```

Best parameters: {'n_estimators': 100, 'max_depth': 5}
Best score: 0.924594961178063

Training XGBClassifier with best hyperparameters

In [29]:

```
from xgboost import XGBClassifier

model = XGBClassifier(max_depth=5, n_estimators=100, class_weight='balanced', n_jobs=-1)
model.fit(train, y_train)
y_probs_train = model.predict_proba(train)
y_probs_test = model.predict_proba(test)
y_predicted_train = model.predict(train)
y_predicted_test = model.predict(test)

# keep probabilities for the positive outcome only
y_probs_train = y_probs_train[:, 1]
y_probs_test = y_probs_test[:, 1]


# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['XGBoost'] = [test_auc, test_acc]

# calculate roc curve
```

```
# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()
```

```
*****
Train AUC: 0.942
Test AUC: 0.854
*****
Train Accuracy: 0.920
Test Accuracy: 0.785
```



Train Confusion Matrix

In [30]:

```
from sklearn.metrics import confusion_matrix

cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```


Confusion matrix:

```
[[31231  739]
 [ 2151 2047]]
```

Out[30]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3acc6eebe0>





Test Confusion Matrix

In [31]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[6284 1668]
 [ 275  816]]
```

Out[31]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3acc446be0>



Stacking Classifier

Hyperparameter tuning meta-classifier (Logistic Regression)

In [34]:

```
from sklearn.calibration import CalibratedClassifierCV
from mlxtend.classifier import StackingClassifier

model_1 = LogisticRegression(C=3, class_weight='balanced', n_jobs=-1)
model_1.fit(train, y_train)
clf_1 = CalibratedClassifierCV(model_1, method='sigmoid')

model_2 = RandomForestClassifier(n_estimators=250, max_depth=25, class_weight='balanced', n_jobs=-1)
```

```

model_2.fit(train, y_train)
clf_2 = CalibratedClassifierCV(model_2, method='sigmoid')

model_3 = SGDClassifier(alpha=0.0001, class_weight='balanced', n_jobs=-1)
model_3.fit(train, y_train)
clf_3 = CalibratedClassifierCV(model_3, method='sigmoid')

model_4 = XGBClassifier(max_depth=5, n_estimators=100, class_weight='balanced', n_jobs=-1)
model_4.fit(train, y_train)
clf_4 = CalibratedClassifierCV(model_4, method='sigmoid')

C = [0.0001, 0.001, 0.01, 0.1, 1, 10]
roc = 0
best_C = 0
for i in C:
    log_reg = LogisticRegression(C=i, n_jobs=-1)
    model = StackingClassifier(classifiers=[clf_1, clf_2, clf_3, clf_4], meta_classifier=log_reg, use_probabilities=True)
    model.fit(train, y_train)
    model_roc = roc_auc_score(y_test, model.predict_proba(test)[:, 1])
    if roc < model_roc:
        roc = model_roc
        best_C = i

```

In [35]:

```
best_C
```

Out[35]:

```
0.0001
```

Training stacking classifier with best hyperparameter for meta-classifier

In [36]:

```

from mlxtend.classifier import StackingClassifier

log_reg = LogisticRegression(C=0.0001, n_jobs=-1)
stack_clf = StackingClassifier(classifiers=[clf_1, clf_2, clf_3, clf_4], meta_classifier=log_reg, use_probabilities=True)
stack_clf.fit(train, y_train)

y_probs_train = stack_clf.predict_proba(train)
y_probs_test = stack_clf.predict_proba(test)
y_predicted_train = stack_clf.predict(train)
y_predicted_test = stack_clf.predict(test)

# keep probabilities for the positive outcome only
y_probs_train = y_probs_train[:, 1]
y_probs_test = y_probs_test[:, 1]

# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['Stacking Classifier'] = [test_auc, test_acc]


# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the stack_clf
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')

```



```
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()
```

```
*****
Train AUC: 0.983
Test AUC: 0.871
*****
Train Accuracy: 0.884
Test Accuracy: 0.879
```



Train Confusion Matrix

In [37]:


```
from sklearn.metrics import confusion_matrix

cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Confusion matrix:
[[31970    0]
 [ 4198    0]]
```

Out[37]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3abc696fd0>





Test Confusion Matrix

In [38]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[7952   0]
 [1091   0]]
```

Out[38]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3ab5f84978>



Voting Classifier

In [40]:

```
from sklearn.ensemble import VotingClassifier
model = VotingClassifier(estimators=[('log_reg', clf_1), ('rf', model_2), ('stack', stack_clf),
('xgb', model_4), ('log_reg_1', model_1)], voting='soft')
model.fit(train, y_train)

y_probs_train = model.predict_proba(train)
y_probs_test = model.predict_proba(test)
y_predicted_train = model.predict(train)
y_predicted_test = model.predict(test)

# keep probabilities for the positive outcome only
y_probs_train = y_probs_train[:, 1]
y_probs_test = y_probs_test[:, 1]

# calculate AUC and Accuracy
train_auc = roc_auc_score(y_train, y_probs_train)
test_auc = roc_auc_score(y_test, y_probs_test)
train_acc = accuracy_score(y_train, y_predicted_train)
test_acc = accuracy_score(y_test, y_predicted_test)
```

```

print('*'*50)
print('Train AUC: %.3f' % train_auc)
print('Test AUC: %.3f' % test_auc)
print('*'*50)
print('Train Accuracy: %.3f' % train_acc)
print('Test Accuracy: %.3f' % test_acc)

score['Voting Classifier'] = [test_auc, test_acc]


# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_probs_train)
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_probs_test)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr, marker='.', label='Train AUC')
plt.plot(test_fpr, test_tpr, marker='.', label='Test AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.show()

```

```

*****
Train AUC: 0.977
Test AUC: 0.892
*****
Train Accuracy: 0.929
Test Accuracy: 0.894

```



Train Confusion Matrix

In [41]:

```

from sklearn.metrics import confusion_matrix

cma = confusion_matrix(y_train, y_predicted_train)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')

```

```


Confusion matrix:
[[31657  313]
 [ 2253 1945]]

```

Out[41]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3abc1dcc50>





Test Confusion Matrix

In [42]:

```
from sklearn.metrics import confusion_matrix


cma = confusion_matrix(y_test, y_predicted_test)
print('Confusion matrix:\n', cma)
df_cm = pd.DataFrame(cma, range(2), columns=range(2))
plt.figure(figsize = (10,7))
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Confusion matrix:

```
[[7585  367]
 [ 595  496]]
```

Out[42]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3ac1115e10>



Conclusion

- It was a great learning experience working on a financial dataset.
- Our dataset consist of categorical and numerical features.
- We have 16 independent features, out of these only half of them are important.

- 'duration' is the most important feature while 'education' is the least important feature.
- Month of May have seen the highest number of clients contacted but have the least success rate. Highest success rate is observed for end month of the financial year as well as the calendar year. So one can say that our dataset have some kind of seasonality.
- When visualized age in groups, it is found that clients with age less than 30 and greater than 60 are less contacted through the campaign but have a higher success rate.
- Different machine learning models are trained and tested on the dataset. Out of those Voting Classifier performs best. Logistic Regression is also an important model as it results in high AUC score.
- Different models are summarized in table below.

In [55]:

```
print('***** Comparison of different models *****')
table = PrettyTable(['Model', 'Test AUC', 'Test Accuracy'])
for item in score.items():
    table.add_row([item[0], item[1][0], item[1][1]])
print(table)
```

```
***** Comparison of different models *****
+-----+-----+-----+
|      Model      |      Test AUC      |      Test Accuracy      |
+-----+-----+-----+
| Logistic Regression | 0.8867956824355852 | 0.7464337056286631 |
|   Random Forest   | 0.7991005150979203 | 0.8122304544951896 |
|         SVM        | 0.8075457787974408 | 0.8053743226805263 |
|       XGBoost      | 0.853617292665249  | 0.7851376755501492 |
| Stacking Classifier | 0.871481409077748  | 0.8793541966161672 |
|   Voting Classifier | 0.8921652047943022 | 0.8936193741015149 |
+-----+-----+-----+
```

References/Citations

1. [Moro et al., 2014] S. Moro, P. Cortez and P. Rita. A Data-Driven Approach to Predict the Success of Bank Telemarketing. Decision Support Systems, Elsevier, 62:22-31, June 2014
2. <https://archive.ics.uci.edu>